

1 Highlights

2 **Sentinel-2 and Landsat 8/9 retrieval of water transparency in Lake** 3 **Titicaca: harmonization limits, validation design and calibrated** 4 **uncertainty**

5 Dina Maribel Yana-Yucra, Richar Andre Vilca-Solorzano, Fred Torres-Cruz,
6 Vladimiro Ibañez

- 7 • Land-derived Landsat/Sentinel-2 harmonization is invalid over dark
8 oligotrophic water
- 9 • Sensors stay 99.8% separable after harmonization; intercepts exceed
10 signal 70-fold
- 11 • Campaign, not station, is the blocking unit in campaign-based moni-
12 toring programmes
- 13 • Honest Secchi retrieval: $R^2=0.574$, RMSE 1.86 m over 812 real match-
14 ups
- 15 • Conformal intervals are calibrated marginally but under-cover in the
16 clearest water

17 Sentinel-2 and Landsat 8/9 retrieval of water
18 transparency in Lake Titicaca: harmonization limits,
19 validation design and calibrated uncertainty

20 Dina Maribel Yana-Yucra^a, Richar Andre Vilca-Solorzano^{a,*}, Fred
21 Torres-Cruz^a, Vladimiro Ibañez^a

*^aFaculty of Statistical and Computer Engineering, Universidad Nacional del Altiplano
(UNAP), Av. Floral 1153, Puno, 21001, Peru*

22 **Abstract**

23 Monitoring water transparency in high-altitude oligotrophic lakes is critical
24 for freshwater management, yet it remains constrained by the sparsity of field
25 data. While multi-mission satellite archives (e.g., harmonized Sentinel-2 and
26 Landsat 8/9) and machine learning offer a synoptic solution, current practices
27 often import assumptions from turbid-water literature that fail in clear-water,
28 campaign-based settings. We assemble 1002 satellite–field match-ups against
29 twelve years of limnological monitoring in Lake Titicaca to address these fail-
30 ures and define a defensible retrieval protocol. First, we demonstrate that
31 standard land-derived harmonization coefficients (e.g., Roy et al., 2016) are
32 invalid over dark water, as their intercepts exceed the native Landsat signal
33 by factors of 1.4 to 69.8, leaving sensors 99.8% separable; we propose a clas-
34 sifier two-sample test to catch this failure. Second, we show that standard
35 station-blocked cross-validation fails to prevent spatial leakage in campaign-
36 based sampling: under campaign-blocked folds, a station-mean null model
37 still explains $R^2 = 0.208$ of the variance. We introduce a null-model diag-

*Corresponding author.

Email addresses: 75923467@est.unap.edu.pe (Dina Maribel Yana-Yucra),
75521963@est.unap.edu.pe (Richar Andre Vilca-Solorzano)

nostic to identify the correct blocking unit (the campaign). Finally, under this strict design, we train a Random Forest regressor to retrieve Secchi disk depth ($R^2 = 0.574$, RMSE 1.86 m) and attach split-conformal prediction intervals. We reveal that while marginally calibrated (89.4% coverage at 90% nominal), conformal intervals structurally under-cover in the clearest water (71.3%). Our findings provide a rigorous blueprint for evaluating machine-learning water quality products in large, oligotrophic lakes.

Keywords: Lake Titicaca, Secchi disk depth, Sentinel-2, Landsat, cross-sensor harmonization, conformal prediction

1. Introduction

Large lakes are sentinels of environmental change, but continuous monitoring of remote, high-altitude systems is constrained by sparse field infrastructure (Williamson et al., 2009). Lake Titicaca (3,810 m a.s.l.), the largest tropical lake in the world by volume, is a critical transboundary freshwater reserve whose water quality is documented only by infrequent field campaigns (OEFA, 2016; Siguayro Mamani et al., 2022), even as eutrophication advances in its near-shore bays (Heredia et al., 2022; Lanza et al., 2024). Satellite remote sensing offers a synoptic solution (Gholizadeh et al., 2016; Topp et al., 2020). Specifically, water transparency—quantified by the Secchi disk depth (Z_{SD})—is an integrative indicator of water clarity strongly linked to light attenuation (Lee et al., 2015). Because Sentinel-2 MSI and Landsat 8/9 OLI provide complementary revisit times, their joint use as a virtual constellation coupled with machine learning is now the standard practice for mapping water quality (Pahlevan et al., 2020; Cao et al., 2022).

62 Although these virtual constellations work well in turbid, productive
 63 “Case-2” waters, they fail to transfer reliably to clear, oligotrophic high-
 64 altitude lakes due to two unresolved methodological assumptions. First,
 65 the cross-sensor harmonization coefficients required to merge Landsat and
 66 Sentinel-2 (e.g., Roy et al., 2016) were fitted on terrestrial targets. Because
 67 terrestrial surfaces are orders of magnitude brighter than clear water, their
 68 additive intercepts overwhelm the over-water signal. While near-infrared
 69 bias over water has been noted (Pahlevan et al., 2019), the magnitude of
 70 this mismatch and its effect on machine-learning retrievals remain unquan-
 71 tified. Second, to prevent spatial autocorrelation from inflating apparent
 72 accuracy, previous methods routinely block cross-validation by monitoring
 73 station (Roberts et al., 2017). However, they assume rather than test that
 74 station identity is the primary source of data leakage. In campaign-based
 75 monitoring programmes, where an entire lake is sampled within a few days
 76 under identical conditions, station-blocking may retain severe leakage and
 77 falsely inflate performance metrics.

78 In this paper, we propose a rigorous validation and retrieval pipeline to
 79 overcome these challenges, using Lake Titicaca as our testing ground. Our
 80 innovation is replacing assumed methodological defaults with data-driven
 81 diagnostics. Specifically, we assemble 1002 real satellite–field match-ups and
 82 apply a classifier two-sample test (adversarial validation) (Lopez-Paz and
 83 Oquab, 2017) to diagnose harmonization failure. Furthermore, we introduce
 84 a null-model diagnostic to identify the true blocking unit for cross-validation.
 85 Finally, we train a Random Forest regressor under the corrected design and
 86 attach split-conformal prediction intervals (Angelopoulos and Bates, 2023)

87 to quantify per-prediction uncertainty.

88 In contrast to previous water-quality mapping efforts (Ayala Izurieta
89 et al., 2023; Maligaya et al., 2024), our study makes three explicit contri-
90 butions:

- 91 **1. Quantification of Harmonization Limits:** We show that land-
92 derived harmonization intercepts exceed the native Landsat signal by
93 up to 69.8-fold over clear water. Consequently, a classifier can identify
94 the source sensor with 99.8% accuracy from harmonized bands, proving
95 that the model learns a sensor label rather than unified optics.
- 96 **2. Null-Model Validation Design:** We demonstrate that under campaign-
97 blocked folds, a station-mean null model still achieves $R^2 = 0.208$,
98 proving that station-blocking fails to prevent leakage. We establish the
99 campaign as the required blocking unit for this class of data.
- 100 **3. Defensible Retrieval with Conformal Uncertainty:** Under this
101 strict design, our model achieves an honest R^2 of 0.574 (RMSE 1.86 m).
102 We further reveal that while split-conformal intervals are marginally
103 calibrated (89.4% coverage), they structurally under-cover extreme val-
104 ues in the clearest water, highlighting a critical limitation for opera-
105 tional monitoring.

106 2. Materials and methods

107 2.1. Study area

108 Lake Titicaca comprises three limnologically distinct zones (Fig. 1): the
109 large, deep, oligotrophic Lago Mayor; the shallower Lago Menor (Wiñay-
110 marca) to the south-east; and Bahía de Puno, which receives municipal and

111 agricultural inputs from the city of Puno and is the most eutrophic sector.
112 Measured transparency spans 1.5 m in near-shore turbid water to 16.5 m in
113 the open oligotrophic basin.

114 2.2. *In-situ data*


115 We used 881 field measurements of Secchi transparency, with co-measured
116 chlorophyll-*a*, suspended solids, temperature, pH, dissolved oxygen and nu-
117 trients, from the limnological monitoring campaigns of IMARPE (Instituto
118 del Mar del Perú) on Lake Titicaca, spanning 2011–2024 across 160 sta-
119 tion identifiers distributed over the whole lake. Six of those identifiers are
120 a 2017 naming variant of a station already recorded under another label,
121 so the records describe 154 distinct stations; we retain the identifiers as
122 issued, which makes station-blocked validation marginally more permissive
123 and therefore does not favour the argument of Section 2.6. Of the 881 mea-
124 surements, 734 include a valid Secchi reading. The July 2019 campaign (71
125 stations) is published by Siguyayro Mamani and Franco (2022); the remaining
126 years derive from the same IMARPE programme. As a provenance check, our
127 2019 zone-level means reproduce the published values almost exactly (Lago
128 Mayor conductivity 1513.4 vs. 1513.5 $\mu\text{S cm}^{-1}$; Lago Menor chlorophyll-*a*
129 1.52 vs. 1.52 mg m^{-3}).

130 Three properties of this record constrain everything that follows, and we
131 state them here rather than in the limitations (Fig. A.5).

132 First, the monitoring programme is itself almost entirely a dry-season
133 programme. The 2011–2024 record contains 13 campaigns with Secchi read-
134 ings, of which exactly one — December 2012, 72 measurements — falls in the
135 austral wet season (December–March). That campaign predates both sen-

(a) Limnological monitoring network (IMARPE)

IMARPE/ALT stations with satellite-field match-ups (2013–2024), coloured by mean measured Secchi depth. The trophic gradient is visible in the field data alone



(b) Sampling effort is very uneven

Each point is a station. 143 stations, but most were visited in very few campaigns, limiting the power of any station-based analysis




Figure 1: Study area and monitoring network. (a) IMARPE/ALT stations with a satellite–field match-up, coloured by mean measured Secchi depth over the Natural Earth lake boundary; the trophic gradient is evident in the field data alone. (b) Sampling effort per station: most stations were visited in few campaigns, which limits the power of any station-level analysis.

sors used here: Landsat 8 became operational in May 2013 and Sentinel-2A in June 2015. Consequently every match-up in our dataset falls between July and October, and this is a property of the field programme rather than of our processing: no choice of match-up window or cloud threshold can recover wet-season pairs that have no contemporaneous imagery. We return to what this means for operational use in Section 4.3.

Second, sampling is episodic: campaigns exist for 2013–2024 with no campaigns in 2020, 2021, 2023.

Third, the reference measurement is coarse: 73 distinct Secchi values occur in 812 match-ups, because disk readings are recorded to the nearest half metre.

No synthetic or spectrally derived labels were used at any stage.

2.3. Satellite data and cross-sensor harmonization

Sentinel-2 MSI (harmonized surface reflectance) and Landsat 8/9 OLI (Collection 2, Level-2; Vermote et al., 2016) were accessed through Google Earth Engine (Gorelick et al., 2017). Clouds were masked with the Sentinel-2 Scene Classification Layer plus cloud probability, and with the Landsat QA_PIXEL bitmask (dilated cloud, cloud and shadow). The cirrus flag was not applied, as it is near-permanently set at 3810 m and would mask the entire archive. We retained the six bands common to both sensors (blue, green, red, NIR, SWIR1, SWIR2).

Landsat surface reflectance was first harmonized to Sentinel-2 equivalent reflectance with the per-band linear coefficients of Roy et al. (2016), of the form $\rho_{\text{MSI}} = a + b \rho_{\text{OLI}}$. Section 3.1 evaluates whether this transform is valid over the water of Lake Titicaca and derives an alternative. For that evalu-

161 ation we identified every in-situ event observed by both sensors within the
162 match-up window — 270 paired station–date observations — and regressed
163 Sentinel-2 on harmonized Landsat reflectance band by band.

164 *2.4. Satellite–field match-ups*

165 For every in-situ measurement we extracted the cloud-masked median
166 reflectance within a 30 m buffer of the station coordinate, from imagery
167 acquired within ± 10 days. The 30 m buffer matches the coarser Landsat pixel
168 and averages residual geolocation error and sensor noise over several pixels
169 while remaining in open water, away from shoreline mixing. The ± 10 -day
170 window balances the scarce cloud-free image availability at 3810 m against
171 the slow temporal change of transparency in this large, deep lake, and is
172 consistent with windows used in operational Z_{SD} retrieval (Kloiber et al.,
173 2002; Zhang et al., 2022).

174 This yielded 1002 match-ups (340 Sentinel-2, 662 Landsat), of which 812
175 have a Secchi value, across 143 stations and 120 campaign dates. Because
176 both sensors can observe the same field event, 272 match-ups are cross-sensor
177 pairs of a single in-situ measurement; collapsing them leaves 540 independent
178 events, and Section 3.2 verifies that accuracy is unchanged, confirming the
179 result is not an artefact of pseudo-replication.

180 *2.5. Features and model*

181 Twelve sensor-agnostic features were computed from the six common
182 bands: the bands themselves plus NDWI, NDTI and the ratios blue/green,
183 green/blue, red/green and blue/red. We trained a Random Forest regres-
184 sor (Breiman, 2001). Reflectance was scaled with a robust (median/IQR)

185 scaler fitted only on the training partition of each fold. All models were
186 implemented with scikit-learn (Pedregosa et al., 2011).

187 Model choice was settled empirically rather than assumed. We com-
188 pared Random Forest, Extremely Randomized Trees, XGBoost (Chen and
189 Guestrin, 2016) and LightGBM, each with hyperparameters optimized by
190 Optuna (Akiba et al., 2019) inside *nested* cross-validation — tuning on an
191 inner campaign-blocked split of each outer training partition, so no outer fold
192 ever informs its own hyperparameters — and benchmarked the resulting gain
193 against the spread across random seeds. Section 3.4 reports the outcome.

194 2.6. Validation design and null models

195 Rather than a single split, we evaluate a hierarchy of increasingly de-
196 manding designs (Roberts et al., 2017): random K -fold; GroupKFold blocked
197 by station; GroupKFold blocked by *campaign date*; GroupKFold blocked by
198 year; a temporal hold-out; and leave-one-zone-out, in which an entire trophic
199 zone is withheld.

200 Crucially, we do not assume which blocking unit matters. We benchmark
201 each design against null models that predict, respectively, the training-set
202 mean of the zone, of the station and of the campaign date. A blocking unit
203 is doing real work if, and only if, the corresponding null model collapses
204 under it. This diagnostic, rather than convention, determines our primary
205 design.

206 Performance is reported as out-of-fold R^2 , RMSE and MAE, and bench-
207 marked against a classical log-linear blue/green band-ratio algorithm in the
208 standard form $\ln Z_{SD} \sim \ln(\rho_{\text{blue}}/\rho_{\text{green}})$ (Kloiber et al., 2002), with predic-
209 tions bounded to the physically observable range, and against a multiband

210 linear model.

211 2.7. Uncertainty and interpretation

212 We attach a per-prediction interval using split-conformal prediction (An-
213 gelopoulos and Bates, 2023): within each fold, a calibration subset of the
214 training partition yields the $(1 - \alpha)$ quantile of absolute residuals, defining
215 a symmetric interval with finite-sample marginal coverage guarantees (Ro-
216 mano et al., 2019). Because a single nominal level can be met by chance, we
217 report empirical coverage at 50, 80, 90 and 95%, and also coverage condi-
218 tional on measured transparency, since marginal calibration does not imply
219 conditional calibration.

220 The model is interpreted with SHAP (Lundberg and Lee, 2017; Lundberg
221 et al., 2020), computed *out of fold* — each observation explained by the model
222 that did not see it — rather than on a model fitted to the whole dataset, which
223 inflates the apparent importance of memorized features. Trends in annual
224 median in-situ transparency were tested with Mann–Kendall and Sen’s slope
225 (Mann, 1945; Sen, 1968) on the full 2011–2024 in-situ record, not on the
226 match-up subset.

227 3. Results

228 3.1. Land-derived harmonization is invalid over oligotrophic water

229 The Roy et al. (2016) transforms carry additive intercepts of 0.0116–
230 0.0448 in reflectance units. Over Lake Titicaca the native Landsat over-water
231 signal is far smaller than this (Table A.2, Fig. 2a). In the near-infrared the
232 median native reflectance is effectively zero, so the intercept exceeds the

233 signal by a factor of 69.8; in SWIR1 the factor is 23.2, in red 10.6, and even
 234 in blue — the brightest water band — 2.5. The harmonized Landsat bands
 235 are consequently dominated by a constant inherited from a land calibration:
 236 their medians exceed the corresponding Sentinel-2 medians by 1.6-fold in
 237 green and 15.4-fold in the near-infrared.

238 A harmonization that succeeded would make the two sensors statistically
 239 indistinguishable over the same water. They are not. A Random Forest clas-
 240 sifier trained to identify the source sensor from the harmonized features alone
 241 reaches 99.8% accuracy against a 66.3% majority-class baseline (Fig. 2b); re-
 242 stricting it to the visible bands gives 97.9%, and to scale-invariant band ratios
 243 99.5%. This is a classifier two-sample test (Friedman, 2003; Lopez-Paz and
 244 Oquab, 2017): if two samples come from the same distribution, no classifier
 245 can separate them above chance. Sensor identity therefore remains fully en-
 246 coded in the harmonized reflectance, and any model trained on the pooled
 247 archive can read it. The test is not new, but to our knowledge it has not
 248 been used as a harmonization acceptance check; it costs one model fit and it
 249 is decisive.

250 Refitting locally on the 270 station-date events observed by both sensors
 251 gives water-specific coefficients (Table A.3). These restore the spectral sig-
 252 nature almost exactly — the median Landsat/Sentinel-2 ratio falls from 15.4
 253 to about 1.0 across all six bands (Fig. 2d) — but per-observation agreement
 254 remains uneven (Fig. 2c). Green recovers usable agreement ($r = 0.71$), and
 255 NIR and SWIR reach $r \approx 0.6$. Blue does not: $r = 0.26$. This is the most
 256 consequential result of the section, because SHAP attributes 72.6% of the
 257 retrieval to green/blue bands and ratios (Section 3.7). The band on which

258 the retrieval most depends is the band on which the two sensors least agree.

259 3.2. *The campaign, not the station, is the blocking unit*

260 Table 1 reports the retrieval under each validation design alongside the
261 null models (Fig. A.6). Blocking by station changes nothing relative to a ran-
262 dom split (both $R^2 = 0.596$), which under the usual reading would be taken
263 as evidence of no spatial leakage. The null models show why that reading
264 is wrong: under campaign-blocked folds the station-mean null still attains
265 $R^2 = 0.208$ and the zone-mean null $R^2 = 0.287$, whereas the campaign-date-
266 mean null collapses to -0.012 . Dependence in this dataset lives between
267 campaigns, not between stations, because a single campaign samples the
268 whole lake within a few days under one set of atmospheric and limnological
269 conditions. Station-blocked folds leave that structure intact.


270 Under campaign-blocked validation — our primary design — the Random
271 Forest attains $R^2 = 0.574$ and RMSE 1.86 m over 812 match-ups, essentially
272 unbiased (-0.02 m). Blocking by year is stricter still ($R^2 = 0.495$). Collaps-
273 ing cross-sensor pseudo-replicates to 540 independent events leaves accuracy
274 unchanged ($R^2 = 0.571$).

275 3.3. *Retrieval accuracy and benchmarks*

276 Against this design the classical blue/green band ratio attains $R^2 = 0.097$
277 (RMSE 2.71 m) and a multiband linear model $R^2 = 0.323$ (RMSE 2.34 m)
278 (Table A.4, Fig. 3). We report one methodological caution alongside these
279 figures. An unbounded implementation of the same band-ratio algorithm —
280 log-transformed and back-transformed without constraining predictions to
281 the observable range — produces a single prediction of 143.5 m against an


(a) Land intercept exceeds the signal

Roy et al. (2016) coefficients, fitted over vegetation and soil, vs native Landsat reflectance over the lake water. The dashed line marks intercept = signal




(b) Sensors remain distinguishable

Accuracy of a Random Forest predicting the source sensor from already harmonized reflectance. If harmonization worked, it would be at chance



(c) Blue is the band that agrees least between sensors

Sentinel-2 reflectance vs Roy-harmonized Landsat in the 270 events observed by both sensors at the same station and date. The dashed line is 1:1



(d) Local recalibration over water does harmonize bands

Median reflectance over water by sensor. Labels are the Landsat / Sentinel-2 ratio: it goes from 15.4x in NIR to ~1.0x in every band




Figure 2: Why land-derived harmonization fails over clear water. (a) The Roy intercept relative to the native over-water Landsat signal, by band. (b) Accuracy of a classifier predicting the source sensor from harmonized features; under a successful harmonization it would sit at chance. (c) Sentinel-2 versus harmonized Landsat reflectance on the 270 paired station-date events; blue agrees worst. (d) Median spectral signature before and after local water recalibration; labels are the Landsat/Sentinel-2 ratio.

Table 1: Validation design, null models and blocking units (Random Forest, $n = 812$ unless noted). A blocking unit is effective only if the corresponding null model collapses under it.

	Setting	R^2	RMSE (m)
Null models	Zone-mean	0.287	2.40
	Station-mean	0.208	2.53
	Campaign-date-mean	-0.012	2.86
Blocking unit	Random K -fold	0.596	1.81
	By station	0.596	1.81
	By campaign date	0.574	1.86
	By year	0.495	2.02
	One record per event ($n = 540$)	0.571	1.77
Zone extrapolation	Leave out Lago Menor ($n = 133$)	0.247	1.90
	Leave out Bahía de Puno ($n = 88$)	0.036	2.11
	Leave out Lago Mayor ($n = 591$)	-0.047	2.52
Temporal	Train ≤ 2019 , test 2022 & 2024 ($n = 226$)	0.507	2.63


282 observed maximum of 16.5 m, which alone drives R^2 to -2.458 and RMSE to
283 5.29 m (Fig. 3d). Benchmarks of classical algorithms reported without phys-
284 ical bounds should be treated with caution: a single unconstrained extrapo-
285 lation manufactures an arbitrarily large advantage for the machine-learning
286 model.

287 3.4. *Model family and tuning do not change the picture*


288 We compared four learner families, each tuned by Optuna inside nested
289 cross-validation, against target transformations and reduced feature sets (Ta-
290 ble A.5). The best configuration — an equal-weight ensemble of the four
291 tuned models — reached $R^2 = 0.583$, an improvement of 0.009 over the
292 untuned Random Forest. That gain does exceed the spread across random
293 seeds (0.003 in R^2), so it is not noise; but it corresponds to 0.02 m of RMSE
294 on a lake whose reference measurement is recorded to the nearest half metre,
295 and it changes no conclusion in this article. Neither log-transforming the
296 target nor tuning individual learners helped more than that, and restricting
297 the feature set cost accuracy.

298 We therefore retain the untuned Random Forest as the reported model:
299 it is simpler, faster and easier to reproduce, and reporting a tuned ensemble
300 for a 0.009 gain would overstate the precision of the product. The ceiling
301 here is set by the data — a dry-season-only record, a reference measurement
302 quantized at half a metre, and a target whose variance is dominated by a
303 three-level trophic gradient — not by the choice of learner.

(a) Predicted vs measured
Random Forest, campaign-blocked
validation. Dashed line is 1:1




(b) Bias concentrates in the bay
Residuals (predicted - measured) by zone.
Diamond is the mean. Bahía de Puno, the
most eutrophic zone, +1.0 m



• Bahía de Puno • Lago Menor • Lago Mayor

(c) The baseline was misimplemented
Same validation design for all.
The published baseline undid the log without
bounds; a single 143 m prediction ruined its R^2



(d) One impossible value ruins R^2
Classical ratio on log scale.
The 143 m point is the one that produced
the $R^2 = -2.46$




Figure 3: Retrieval under the primary design. (a) Out-of-fold predicted versus measured Secchi depth. (b) Residuals by trophic zone; the retrieval overestimates in eutrophic Bahía de Puno. (c) Benchmark against classical and linear baselines. (d) The unbounded classical implementation in log scale: one impossible prediction destroys the metric.

304 3.5. Where the retrieval stops working

305 Withholding an entire trophic zone collapses the retrieval (Table 1): R^2
306 falls to 0.247 for Lago Menor, 0.036 for Bahía de Puno and -0.047 for Lago
307 Mayor. The model interpolates within optical regimes it has seen and does
308 not extrapolate to a water type absent from training. On the temporal hold-
309 out (train ≤ 2019 , test 2022 and 2024) it retains $R^2 = 0.507$, though 57 of the
310 69 test stations also appear in training (82.6%), so this measures temporal
311 rather than fully independent generalization.

312 Error is unevenly distributed (Table A.6). The retrieval carries a positive
313 bias of 1.01 m in Bahía de Puno, the most eutrophic and most vulnerable
314 zone, because the clear-water majority pulls predictions upward.

315 Two features of the harmonized archive should be read together with these
316 numbers. Removing the two bands most corrupted by the land-derived in-
317 tercepts lowers R^2 only from 0.574 to 0.541; using scale-invariant ratios alone
318 gives 0.468; and applying our water-specific recalibration gives 0.518. Train-
319 ing on each sensor separately gives 0.563 (Sentinel-2, $n = 274$) and 0.549
320 (Landsat, $n = 538$). The pooled model’s small advantage over these vari-
321 ants is consistent with it exploiting sensor identity as an additional covariate
322 rather than with genuine cross-sensor synergy.

323 3.6. Calibrated per-prediction uncertainty

324 Split-conformal intervals are well calibrated *marginally* at every nominal
325 level tested: empirical coverage is 50.9% at 50% nominal, 79.1% at 80%,
326 89.4% at 90% and 94.3% at 95%, a maximum calibration error of 0.94 per-
327 centage points (Fig. A.7a).

They are not calibrated *conditionally*, and this matters more for operational use than the marginal result. Coverage falls to 71.3% in the clearest water ($Z_{SD} > 12.5$ m) and 79.5% in the most turbid ($Z_{SD} < 5$ m), while exceeding the nominal level in the middle of the range (Fig. A.7c). The mechanism is structural rather than incidental: split conformal derives a single global quantile of absolute residuals, so the interval has constant width (6.12 m at 90%) everywhere and cannot adapt to the local difficulty of the retrieval. It therefore buys its marginal guarantee by over-covering in the well-sampled middle of the range and under-covering at both extremes — which are precisely where management decisions are made. A locally adaptive variant such as conformalized quantile regression (Romano et al., 2019) would be required for reliable conditional coverage, and we recommend reporting coverage conditional on the predicted value as standard practice whenever conformal intervals accompany an operational product.

For a lake whose transparency spans 1.5–16.5 m, a ± 3 m interval is informative for zone-scale contrasts and multi-metre change, and is not informative for sub-metre differences.

3.7. Interpretation and retrievability

Out-of-fold SHAP attributes 72.6% of the retrieval to green and blue bands and their ratios, 14.4% to NIR and SWIR, and the remainder to other indices (Fig. 4a). Green/blue alone accounts for 37%. The dependence is monotone and of the sign expected from water-clarity optics (Fig. 4b): a higher green-to-blue ratio pushes the prediction toward less transparent water. The retrieval is therefore grounded in optics rather than location memorization — while resting on the band pair with the worst cross-sensor

353 agreement.

354 With real labels and the same campaign-blocked design, no other vari-
355 able is retrievable (Fig. 4c): chlorophyll-*a* reaches $R^2 = 0.012$, suspended
356 solids 0.056, in-situ water temperature 0.113, and Landsat thermal surface
357 temperature 0.020 ($r = 0.14$, $n = 662$). Their in-situ ranges — chlorophyll-*a*
358 0–4.6 mg m⁻³, suspended solids mostly below 10 mg L⁻¹ — lie below the op-
359 tical and thermal detection limits of these sensors. We subject the negative
360 claim to the same design as the positive one, and report it to caution against
361 applying chlorophyll algorithms (Gitelson et al., 2009; Mishra and Mishra,
362 2012) or thermal retrievals (Ruiz-Verdú et al., 2016) in clear high-altitude
363 water without local validation.


364 3.8. Temporal patterns and the fragility of trend claims

365 Mann–Kendall tests on annual median in-situ transparency over the full
366 2011–2024 record show no statistically significant trend in any zone (Bahía
367 de Puno: Sen slope 0.16 m yr⁻¹, $p = 0.241$; Lago Mayor: 0.25 m yr⁻¹,
368 $p = 0.138$; Fig. A.8a).

369 That conclusion is fragile in a way worth reporting explicitly. Restricting
370 the series to 2013 onwards — the natural choice if trends were computed
371 on the satellite match-up subset, whose record begins with Landsat 8 —
372 turns both Bahía de Puno ($p = 0.017$) and Lago Mayor (Sen +0.42 m yr⁻¹,
373 $p = 0.012$) significantly positive. Across start years from 2011 to 2014, 5 of
374 12 zone–start-year combinations change verdict (Fig. A.8b). The year 2013
375 has the lowest annual medians on record, so beginning there manufactures
376 an upward trend. With 9–11 annual values and three gap years the test has
377 low power, and the correct reading is absence of evidence, not evidence of


(a) Green and blue dominate attribution

Mean |SHAP| computed OUT-OF-FOLD: each observation is explained by the model that did not see it. Labels are the share of the total



(b) How Green/Blue acts

SHAP value vs. feature value. Greener water relative to blue pushes the prediction toward less transparent water. Trimmed to the 1-99th percentile (18 extreme points out)



(c) Only transparency is retrievable in this water

Random Forest with the same honest design (campaign blocking) for each variable. The negative claim is subject to the same rigor as the positive one




Figure 4: Interpretation and operational envelope. (a) Out-of-fold SHAP importance; green and blue dominate. (b) SHAP dependence on the leading feature, by trophic zone. (c) Retrievability of each water-quality variable under the same campaign-blocked design; only transparency clears a useful threshold.

378 stability. We recommend that trend analyses in campaign-based lake records
379 report this sensitivity as a matter of course.

380 4. Discussion

381 4.1. Implications for harmonized multi-mission archives

382 The virtual-constellation approach (Claverie et al., 2018; Pahlevan et al.,
383 2019) has become the default for multi-decadal lake monitoring, and the Roy
384 et al. (2016) coefficients are its most reused transform. Our results place
385 a boundary on where that default is safe. Those coefficients were fitted on
386 terrestrial targets and carry additive intercepts appropriate to surfaces one
387 to two orders of magnitude brighter than clear water. Over Lake Titicaca the
388 intercept exceeds the native over-water signal by a factor of 69.8 in the near-
389 infrared and by 2.5 even in blue, and the two sensors remain 99.8% separable
390 after the transform. That water reflectance from these products is unreliable
391 is already documented (Pahlevan et al., 2019; Masek et al., 2021); what our
392 paired observations add is the magnitude — the problem is not confined
393 to the near-infrared, and it is large enough to make the harmonized bands
394 function as sensor labels for a machine-learning model — and the finding that
395 local recalibration repairs the aggregate spectral signature without repairing
396 per-observation agreement in blue.

397 The practical consequence is direct. A machine-learning model trained
398 on a pooled, land-harmonized archive can achieve apparently good accuracy
399 while having learned two sensor-specific mappings joined by a readable sensor
400 label. That is not a harmonized retrieval, and it will not behave as one when
401 applied to a scene from either sensor alone. The diagnostic is cheap and we

402 recommend it as routine: train a classifier to predict the source sensor from
403 the harmonized features and report its accuracy against the majority-class
404 baseline. If the sensors remain separable, the harmonization has not done
405 its job. For clear oligotrophic lakes we further recommend deriving water-
406 specific coefficients from same-day paired observations, reporting per-band
407 agreement, and — where blue-band agreement is as poor as we find here —
408 reporting single-sensor retrievals alongside any pooled product.

409 *4.2. Implications for validation design*

410 Blocking by station is the standard remedy for spatial autocorrelation
411 in this literature, and in our data it changes nothing: random and station-
412 blocked folds give identical accuracy. Read conventionally, that would be
413 reassuring. The null models show it is not. Under campaign-blocked folds
414 the station-mean null still explains 21% of variance and the zone-mean null
415 29%, while the campaign-date-mean null explains none. In a campaign-based
416 monitoring programme the whole lake is sampled within a few days under
417 shared atmospheric and limnological conditions, so observations cluster by
418 campaign far more strongly than by location; station-blocked folds leave that
419 structure entirely intact.

420 The general point is that the blocking unit should be diagnosed, not
421 assumed. Reporting null models under each candidate design costs little
422 and identifies the unit that carries the dependence. This applies well beyond
423 Lake Titicaca: most large-lake monitoring programmes sample in campaigns,
424 so the same mismatch between the conventional blocking unit and the real
425 dependence structure is likely to be widespread.

426 For our retrieval, the honest figure is $R^2 = 0.574$ (RMSE 1.86 m) under

427 campaign-blocked validation, against 0.287 for a zone-mean null: the model
428 roughly doubles the skill available from knowing only which basin a sample
429 came from. That is a real but modest contribution, and stating it that way
430 is more useful to a manager than an unqualified R^2 .

431 4.3. Operational envelope for Lake Titicaca

432 Our RMSE of 1.86 m spans roughly 12% of the observed Z_{SD} range
433 and sits below the published range for Z_{SD} retrieval ($R^2 \approx 0.6$ – 0.9 ; Zhang
434 et al., 2022; Lee et al., 2015), obtained here under a predominantly clear
435 oligotrophic regime and a stricter validation design. It is also approaching
436 the granularity of the reference measurement itself: disk readings are recorded
437 to the nearest half metre, giving only 73 distinct values across 812 match-ups.

438 Four constraints bound operational use. The retrieval carries a 1.01 m
439 positive bias in Bahía de Puno — the most eutrophic and most vulnerable
440 zone — so mapped values there must be read with their intervals and con-
441 firmed by field sampling. Leave-one-zone-out results show the model cannot
442 extrapolate to an unsampled optical regime; integration must be incremen-
443 tal, each new campaign both validating the current map and extending the
444 training set. The retrieval is calibrated and validated for the austral dry sea-
445 son only, and the reason is worth stating precisely because it determines the
446 remedy. It is not that wet-season imagery is unusable at 3810 m, nor that our
447 match-up criteria excluded it: it is that the field programme has produced
448 a single wet-season campaign in fourteen years, and that campaign predates
449 both satellites. Wet-season blooms in Bahía de Puno — the phenomenon a
450 transparency product would most usefully warn about — are therefore en-
451 tirely unrepresented, and the model should not be applied between December

452 and March.

453 This converts the limitation into a specific, actionable recommendation
454 for the monitoring authorities. Extending the protocol to the wet season does
455 not require a better algorithm; it requires field data that do not yet exist. A
456 single additional wet-season campaign per year, timed to any cloud-free win-
457 dow and covering the three trophic zones, would within a few years supply
458 enough independent campaign-blocks to calibrate and validate a wet-season
459 retrieval under the same design used here. Because the binding constraint is
460 sampling rather than sensing, this is among the cheapest available improve-
461 ments to Lake Titicaca’s water-quality monitoring capability. And because
462 conformal coverage degrades at both extremes of the transparency range, the
463 nominal 90% interval should be treated as optimistic in the clearest and most
464 turbid water.

465 Within those bounds the product is useful to the regional monitoring
466 authorities (IMARPE, ALT, OEFA, and the binational management frame-
467 work). A calibrated ± 3 m interval resolves zone-scale contrasts and multi-
468 metre change, which is the scale at which clarity loss in Bahía de Puno
469 would be actionable, and a single transboundary product under one cali-
470 bration supports comparable assessment across the Peru–Bolivia boundary
471 where coordinated field monitoring is logistically difficult. It does not resolve
472 sub-metre change and should not be presented as doing so.

473 4.4. *Limitations*

474 Beyond the constraints already stated, three limitations bound this work.
475 The in-situ record is episodic, with no campaigns in 2020, 2021, 2023 and only
476 one wet-season campaign in fourteen years, so the temporal hold-out rests

477 on two test years, the trend analysis on 9–11 annual values, and the seasonal
478 envelope on none. Transparency above approximately 15 m saturates the
479 optical signal. And this is a single lake: the two diagnostics we propose are
480 general, but the water-specific coefficients in Table A.3 are local to Titicaca
481 and must be refitted elsewhere.

482 5. Conclusions

483 Using 1002 satellite–field match-ups over twelve years, we show that two
484 practices imported from the turbid-water remote-sensing literature fail in a
485 clear high-altitude large lake. Cross-sensor harmonization coefficients derived
486 over land leave Sentinel-2 and Landsat 99.8% separable over oligotrophic
487 water, because their additive intercepts exceed the native over-water signal
488 by up to a factor of 69.8; and validation blocked by monitoring station leaves
489 the dominant dependence — between campaigns — entirely intact.

490 Correcting both, a Random Forest retrieves Secchi depth with $R^2 = 0.574$
491 and RMSE 1.86 m, roughly double the skill of a zone-mean null, with confor-
492 mal prediction intervals that are calibrated marginally but under-cover at the
493 extremes of the transparency range. Skill collapses under extrapolation to an
494 unsampled trophic zone; chlorophyll-*a*, suspended solids and thermal tem-
495 perature are not retrievable at all; and no transparency trend is significant
496 once the full in-situ record is used rather than the satellite-era subset.

497 We therefore recommend two routine checks for multi-mission lake water-
498 quality products, neither of which requires new methodology: report the
499 accuracy of a classifier two-sample test on the harmonized features, and re-
500 port null models under each candidate validation blocking unit. Both cost

one model fit, and both distinguish genuine optical retrieval from sensor labelling and sampling design. For Lake Titicaca the resulting product is a dry-season, zone-scale transparency indicator with honest ± 3 m intervals — less than has sometimes been claimed for such lakes, and more than the region has had.

Appendix A. Supplementary Figures and Tables

Table A.2: Effect of the Roy et al. (2016) OLI→MSI transform over Lake Titicaca water. Native OLI is the Collection-2 surface reflectance recovered by inverting the transform. Reflectance is dimensionless.


Band	Roy intercept a	Native OLI median	Intercept / signal	Harmonized LS median	Harmonized / S2 median
Blue	0.0183	+0.0072	2.5	0.0247	1.74
Green	0.0123	+0.0089	1.4	0.0206	1.61
Red	0.0123	+0.0013	10.6	0.0134	2.66
NIR	0.0448	−0.0006	69.8	0.0443	15.42
SWIR1	0.0306	+0.0013	23.2	0.0317	13.34
SWIR2	0.0116	+0.0013	8.2	0.0129	5.70

CRedit authorship contribution statement

Dina Maribel Yana-Yucra: Conceptualization, Methodology, Investigation, Data curation, Writing – original draft. **Richar Andre Vilca-Solorzano:** Methodology, Software, Formal analysis, Validation, Visualiza-


(a) Three years with no campaigns and no data outside the dry season

Number of satellite-field match-ups by year and month. There is no data for 2020, 2021 or 2023, so the 2022-2024 hold-out only contains 2022 and 2024




(b) Only dry season data exists

Bahía de Puno blooms occur during the wet season, which the dataset does not cover at all



(c) The trophic gradient dominates Secchi

Only 73 distinct values in 812 match-ups: readings are rounded to half a metre, so the RMSE of 1.86 m borders the granularity of the reference measurement itself



■ Wet season (Dec-Mar) ■ Transition ■ Dry season (Jul-Oct)

Figure A.5: What the dataset actually contains. (a) Match-ups by year and month; three years have no campaign at all, so a hold-out described as “2022–2024” contains only 2022 and 2024. (b) Seasonal coverage: zero match-ups in the wet season, when blooms occur in Bahía de Puno. (c) Secchi distribution by trophic zone; the reference measurement is quantized at roughly half a metre.


Table A.3: Water-specific LS→S2 recalibration fitted on 270 paired station–date observations, of the form $\rho_{S2} = a + b \rho_{LS, Roy}$.

Band	a	b	Pearson r	Agreement
Blue	+0.0084	0.258	0.26	unusable
Green	−0.0009	0.665	0.71	good
Red	−0.0005	0.427	0.52	weak
NIR	−0.0250	0.631	0.62	acceptable
SWIR1	−0.0223	0.770	0.64	acceptable
SWIR2	−0.0057	0.615	0.60	acceptable

(a) Only campaign blocking neutralizes its own null model

R^2 of models predicting only the group mean, with folds blocked by campaign.

A blocking unit is effective if and only if its own null model collapses to zero under it



(b) Station blocking changes nothing; extrapolation breaks everything

Out-of-fold R^2 of the Random Forest; the large point is the primary design.

The fact that station blocking matches with random does not prove absence of leakage: it proves that it retains nothing






Figure A.6: Diagnosing the blocking unit. (a) Null models evaluated under campaign-blocked folds: only the campaign-mean null collapses, so only campaign blocking is doing work. (b) Out-of-fold R^2 under each validation design and under extrapolation to an unseen trophic zone. That station blocking matches the random split is not evidence of no leakage; it is evidence that station blocking retains nothing.

(a) Calibrated, on average
Out-of-fold empirical coverage vs nominal. Gray band is ± 3 points



(c) But not at the ends of the range
Split-conformal uses a single global quantile, so the width is constant (~ 6.1 m) and the interval cannot adapt: it undercovers in the clearest water



(b) The 90% interval vs the measurement
Match-ups ordered by measured transparency. The line is the prediction, the band is the 90% conformal interval. 86 out of 812 points fall outside





Figure A.7: Conformal prediction intervals under campaign-blocked validation. (a) Empirical versus nominal coverage: marginally calibrated across the 50–95% range. (b) The 90% interval against the measured value. (c) Coverage conditional on measured transparency. Because split conformal uses a single global residual quantile, the interval width is constant (≈ 6.12 m) and cannot adapt, so the marginal guarantee is bought by over-covering in the middle of the range and under-covering at both extremes.

(a) No zone shows significant trend in the full record

Annual median transparency in-situ (2011–2024, 734 readings). The vertical dotted lines mark the years with no campaign: 2020, 2021 and 2023



(b) The verdict is set by the start year, not the lake

Same Mann-Kendall test, same data, different start year. Starting in 2013 — the year with the lowest medians in the record — turns two zones into "significantly more transparent"




Figure A.8: Transparency trends and their fragility. (a) Annual median in-situ transparency by zone over the full 2011–2024 record, with Mann–Kendall/Sen statistics; no trend is significant. (b) The same test with different start years: beginning the series in 2013, the lowest-median year on record, turns two zones significantly positive.

Table A.4: Head-to-head benchmark, campaign-blocked GroupKFold, $n = 812$.

Algorithm	R^2	RMSE	MAE	Bias
Classical blue/green, unbounded ^a	−2.458	5.29	1.91	−0.19
Classical blue/green, bounded ^b	0.097	2.71	2.13	−0.46
Multiband linear	0.323	2.34	1.78	−0.01
Random Forest (this study)	0.574	1.86	1.44	−0.02

RMSE, MAE and bias in metres.

^a back-transformed without bounding predictions; see text.

^b standard form of Kloiber et al. (2002).

511 tion, Writing – original draft, Writing – review & editing. **Fred Torres-**
512 **Cruz:** Supervision, Resources, Writing – review & editing. **Vladimiro**
513 **Ibañez:** Supervision, Validation, Writing – review & editing.

514 Declaration of competing interest

515 The authors declare that they have no known competing financial inter-
516 ests or personal relationships that could have appeared to influence the work
517 reported in this paper.

Table A.5: Model selection under nested cross-validation (Optuna inside an inner campaign-blocked split of each outer training partition). Selected configurations; the full grid is in the repository.

Model	Features	R^2	RMSE (m)
Ensemble of the four tuned models	all 12	0.583	1.84
LightGBM, tuned	all 12	0.581	1.84
XGBoost, tuned	all 12	0.577	1.85
Random Forest, tuned	all 12	0.576	1.85
Random Forest, untuned (reported)	all 12	0.574	1.86
Extremely Randomized Trees, untuned	all 12	0.573	1.86
XGBoost, tuned, $\log Z_{SD}$	all 12	0.561	1.89
LightGBM, tuned	visible only	0.544	1.92
LightGBM, tuned	ratios only	0.488	2.04
Seed-to-seed spread of the untuned Random Forest: R^2 range 0.003 over five seeds.			

Table A.6: Out-of-fold error by trophic zone, campaign-blocked validation.

Zone	R^2	RMSE (m)	Bias (m)	n
Bahía de Puno	0.221	1.90	+1.01	88
Lago Menor	0.211	1.94	+0.08	133
Lago Mayor	0.448	1.83	−0.19	591

518 **Data availability**

The satellite–field match-up datasets and the complete analysis pipeline, including the scripts that generate every number, table and figure in this article, are openly available at:

519 <https://github.com/Andre031222/harmonized-titicaca-transparency>



Repository

Raw imagery is freely available via Google Earth Engine; the in-situ measurements derive from published IMARPE/ALT and OEFA monitoring records.

520 **Acknowledgements**

521 We thank IMARPE, ALT and OEFA for the publicly available monitoring
522 records that made this study possible. This research received no external
523 funding.

524 **References**

- 525 Akiba, T., Sano, S., Yanase, T., Ohta, T., Koyama, M., 2019. Optuna: A
526 next-generation hyperparameter optimization framework, in: Proceedings
527 of the 25th ACM SIGKDD International Conference on Knowledge Dis-
528 covery and Data Mining, pp. 2623–2631. doi:10.1145/3292500.3330701.
- 529 Angelopoulos, A.N., Bates, S., 2023. Conformal prediction: A gentle in-
530 troduction. Foundations and Trends in Machine Learning 16, 494–591.
531 doi:10.1561/22000000101.

532 Ayala Izurieta, J.E., Beltrán Dávalos, A.A., Jara Santillán, C.A.,
 533 Godoy Ponce, S.C., Van Wittenberghe, S., Verrelst, J., Delegido, J.,
 534 2023. Spatial and temporal analysis of water quality in high Andean
 535 lakes with Sentinel-2 satellite automatic water products. *Sensors* 23, 8774.
 536 doi:10.3390/s23218774.

537 Breiman, L., 2001. Random forests. *Machine Learning* 45, 5–32. doi:10.102
 538 3/A:1010933404324.

539 Cao, Z., Ma, R., Liu, M., Duan, H., Xiao, Q., Xue, K., Shen, M., 2022.
 540 Harmonized chlorophyll-a retrievals in inland lakes from Landsat-8/9 and
 541 Sentinel 2A/B virtual constellation through machine learning. *IEEE Trans-*
 542 *actions on Geoscience and Remote Sensing* 60, 1–16. doi:10.1109/TGRS.2
 543 022.3207345.

544 Chen, T., Guestrin, C., 2016. XGBoost: A scalable tree boosting system,
 545 in: *Proceedings of the 22nd ACM SIGKDD International Conference on*
 546 *Knowledge Discovery and Data Mining*, pp. 785–794. doi:10.1145/2939
 547 672.2939785.

548 Claverie, M., Ju, J., Masek, J.G., Dungan, J.L., Vermote, E.F., Roger, J.C.,
 549 Skakun, S.V., Justice, C., 2018. The harmonized landsat and Sentinel-2
 550 surface reflectance data set. *Remote Sensing of Environment* 219, 145–161.
 551 doi:10.1016/j.rse.2018.09.002.

552 Friedman, J.H., 2003. On multivariate goodness-of-fit and two-sample test-
 553 ing. *Statistical Problems in Particle Physics, Astrophysics and Cosmology*
 554 (PHYSTAT) , 311–313.

555 Gholizadeh, M.H., Melesse, A.M., Reddi, L., 2016. A comprehensive review
556 on water quality parameters estimation using remote sensing techniques.
557 Sensors 16, 1298. doi:10.3390/s16081298.

558 Gitelson, A.A., Gurlin, D., Moses, W.J., Barrow, T., 2009. A bio-optical
559 algorithm for the remote estimation of the chlorophyll-a concentration in
560 case 2 waters. Environmental Research Letters 4, 045003. doi:10.1088/
561 1748-9326/4/4/045003.

562 Gorelick, N., Hancher, M., Dixon, M., Ilyushchenko, S., Thau, D., Moore,
563 R., 2017. Google Earth Engine: Planetary-scale geospatial analysis for
564 everyone. Remote Sensing of Environment 202, 18–27. doi:10.1016/j.rs
565 e.2017.06.031.

566 Heredia, C., Guédron, S., Point, D., Perrot, V., Campillo, S., Verin, C., Es-
567 pinoza, M.E., Fernandez, P., Duwig, C., Achá, D., 2022. Anthropogenic
568 eutrophication of Lake Titicaca (Bolivia) revealed by carbon and nitro-
569 gen stable isotopes fingerprinting. Science of the Total Environment 845,
570 157286. doi:10.1016/j.scitotenv.2022.157286.

571 Kloiber, S.M., Brezonik, P.L., Olmanson, L.G., Bauer, M.E., 2002. A proce-
572 dure for regional lake water clarity assessment using Landsat multispectral
573 data. Remote Sensing of Environment 82, 38–47. doi:10.1016/S0034-4
574 257(02)00022-6.

575 Lanza, W.G., Hernández, V.C., Achá, D., Lazzaro, X., 2024. Responses
576 of phytoplankton and periphyton community structure to an anthropic

577 eutrophication gradient in tropical high-altitude Lake Titicaca. *Journal of*
578 *Great Lakes Research* 50, 102294. doi:10.1016/j.jglr.2024.102294.

579 Lee, Z., Shang, S., Hu, C., Du, K., Weidemann, A., Hou, W., Lin, J., Lin,
580 G., 2015. Secchi disk depth: A new theory and mechanistic model for
581 underwater visibility. *Remote Sensing of Environment* 169, 139–149. doi:10
582 .1016/j.rse.2015.08.002.

583 Lopez-Paz, D., Oquab, M., 2017. Revisiting classifier two-sample tests, in:
584 International Conference on Learning Representations (ICLR). URL: <https://openreview.net/forum?id=SJkXfE5xx>.
585

586 Lundberg, S.M., Erion, G., Chen, H., DeGrave, A., Prutkin, J.M., Nair,
587 B., Katz, R., Himmelfarb, J., Bansal, N., Lee, S.I., 2020. From local
588 explanations to global understanding with explainable AI for trees. *Nature*
589 *Machine Intelligence* 2, 56–67. doi:10.1038/s42256-019-0138-9.

590 Lundberg, S.M., Lee, S.I., 2017. A unified approach to interpreting model
591 predictions, in: *Advances in Neural Information Processing Systems* 30
592 (NeurIPS 2017), pp. 4765–4774. URL: [https://proceedings.neurips.](https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html)
593 [cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.](https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html)
594 [html](https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html).

595 Maligaya, V.H., Baltodano, A., Agramont, A., van Griensven, A., 2024. Ex-
596 ploring trends and variability of water quality over Lake Titicaca using
597 global remote sensing products. *Remote Sensing* 16, 4785. doi:10.3390/
598 rs16244785.

599 Mann, H.B., 1945. Nonparametric tests against trend. *Econometrica* 13,
600 245–259. doi:10.2307/1907187.

601 Masek, J.G., Ju, J., Roger, J.C., Skakun, S., Vermote, E., Claverie, M.,
602 Dungan, J., Yin, Z., Freitag, B., Justice, C., 2021. HLS Harmonized
603 Landsat Sentinel-2 Product User Guide, Version 2.0. Technical Report.
604 NASA Goddard Space Flight Center / Land Processes DAAC.

605 Mishra, S., Mishra, D.R., 2012. Normalized difference chlorophyll index: A
606 novel model for remote estimation of chlorophyll-a concentration in turbid
607 productive waters. *Remote Sensing of Environment* 117, 394–406. doi:10
608 .1016/j.rse.2011.10.016.

609 OEFA, 2016. Informe 076-2016-OEFA/DE-SDCA: Evaluación de la calidad
610 ambiental de la Bahía Interior de Puno, lago Titicaca. Technical Report.
611 Organismo de Evaluación y Fiscalización Ambiental, Perú.

612 Pahlevan, N., Chittimalli, S.K., Balasubramanian, S.V., Vellucci, V., 2019.
613 Sentinel-2/Landsat-8 product consistency and implications for monitoring
614 aquatic systems. *Remote Sensing of Environment* 220, 19–29. doi:10.101
615 6/j.rse.2018.10.027.

616 Pahlevan, N., Smith, B., Schalles, J., Binding, C., Cao, Z., Ma, R., Alikas,
617 K., Kangro, K., Gurlin, D., Hà, N., Matsushita, B., Moses, W., Greb, S.,
618 Lehmann, M.K., Ondrusek, M., Oppelt, N., Stumpf, R., 2020. Seamless
619 retrievals of chlorophyll-a from Sentinel-2 (MSI) and Sentinel-3 (OLCI) in
620 inland and coastal waters: A machine-learning approach. *Remote Sensing*
621 of *Environment* 240, 111604. doi:10.1016/j.rse.2019.111604.

- 622 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel,
623 O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J.,
624 Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, É., 2011.
625 Scikit-learn: Machine learning in Python. *Journal of Machine Learning*
626 *Research* 12, 2825–2830. URL: [https://jmlr.org/papers/v12/pedreg](https://jmlr.org/papers/v12/pedregosa11a.html)
627 [osa11a.html](https://jmlr.org/papers/v12/pedregosa11a.html).
- 628 Roberts, D.R., Bahn, V., Ciuti, S., Boyce, M.S., Elith, J., Guillerá-Arroita,
629 G., Hauenstein, S., Lahoz-Monfort, J.J., Schröder, B., Thuiller, W.,
630 Warton, D.I., Wintle, B.A., Hartig, F., Dormann, C.F., 2017. Cross-
631 validation strategies for data with temporal, spatial, hierarchical, or phy-
632 logenetic structure. *Ecography* 40, 913–929. doi:10.1111/ecog.02881.
- 633 Romano, Y., Patterson, E., Candès, E.J., 2019. Conformalized quantile
634 regression, in: *Advances in Neural Information Processing Systems* 32
635 (NeurIPS 2019), pp. 3543–3553. URL: [https://proceedings.neurip](https://proceedings.neurips.cc/paper/2019/hash/5103c3584b063c431bd1268e9b5e76fb-Abstract.html)
636 [s.cc/paper/2019/hash/5103c3584b063c431bd1268e9b5e76fb-Abstrac](https://proceedings.neurips.cc/paper/2019/hash/5103c3584b063c431bd1268e9b5e76fb-Abstract.html)
637 [t.html](https://proceedings.neurips.cc/paper/2019/hash/5103c3584b063c431bd1268e9b5e76fb-Abstract.html).
- 638 Roy, D.P., Kovalsky, V., Zhang, H.K., Vermote, E.F., Yan, L., Kumar, S.S.,
639 Egorov, A., 2016. Characterization of Landsat-7 to Landsat-8 reflective
640 wavelength and normalized difference vegetation index continuity. *Remote*
641 *Sensing of Environment* 185, 57–70. doi:10.1016/j.rse.2015.12.024.
- 642 Ruiz-Verdú, A., Jiménez, J.C., Lazzaro, X., Tenjo, C., Delegido, J., Pereira,
643 M., Sobrino, J.A., Moreno, J.F., 2016. Comparison of MODIS and
644 Landsat-8 retrievals of chlorophyll-a and water temperature over Lake Tit-

645 icaca, in: 2016 IEEE International Geoscience and Remote Sensing Sym-
646 posium (IGARSS), pp. 7643–7646. doi:10.1109/IGARSS.2016.7730993.

647 Sen, P.K., 1968. Estimates of the regression coefficient based on Kendall’s
648 tau. *Journal of the American Statistical Association* 63, 1379–1389. doi:10
649 .1080/01621459.1968.10480934.

650 Siguayro Mamani, H., Franco, Y.B., 2022. Aspectos fisicoquímicos del lago
651 Titicaca, julio 2019. *Informe Instituto del Mar del Perú* 49, 538–551.

652 Siguayro Mamani, H., Pasapera, J., Villanueva, C., Coila, Y., Gamarra, C.,
653 2022. Evaluación de fuentes contaminantes en el anillo circunlacustre del
654 lago Titicaca (sector peruano), 2017. *Boletín Instituto del Mar del Perú*
655 37, 361–386. doi:10.53554/boletin.v37i2.378.

656 Topp, S.N., Pavelsky, T.M., Jensen, D., Simard, M., Ross, M.R.V., 2020.
657 Research trends in the use of remote sensing for inland water quality
658 science: Moving towards multidisciplinary applications. *Water* 12, 169.
659 doi:10.3390/w12010169.

660 Vermote, E., Justice, C., Claverie, M., Franch, B., 2016. Preliminary analysis
661 of the performance of the Landsat 8/OLI land surface reflectance product.
662 *Remote Sensing of Environment* 185, 46–56. doi:10.1016/j.rse.2016.0
663 4.008.

664 Williamson, C.E., Saros, J.E., Vincent, W.F., Smol, J.P., 2009. Lakes and
665 reservoirs as sentinels, integrators, and regulators of climate change. *Lim-
666 nology and Oceanography* 54, 2273–2282. doi:10.4319/lo.2009.54.6_p
667 art_2.2273.

668 Zhang, Y., Shi, K., Sun, X., Zhang, Y., Li, N., Wang, W., Zhou, Y., Zhi,
669 W., Liu, M., Li, Y., 2022. Improving remote sensing estimation of Secchi
670 disk depth for global lakes and reservoirs using machine learning methods.
671 GIScience & Remote Sensing 59, 1367–1383. doi:10.1080/15481603.202
672 2.2116102.